


T = 1 s · 4 release slots

Illustrative schedule



Capacity → prefetch → reuse

300 ms: defer

23 + 5 > 25

380 ms: S1 evicts

23 - 6 + 5 = 22

440 → 530 ms

S3 ready → compute

■ GPU KV    □ Host copy    ▨ Copying    ◇ Release    ↑ H2D ↓ D2H    + New KV

G / H: GPU / Host    □ Ready: idle, KV restored    Pool includes H2D targets