


Per-request execution lanes inside the engine — scheduler-trace replay, vanilla vLLM-realtime + parallel ingest, N=8

One tick fully resolved (cell = one engine step for that session; edges = real scheduler timestamps)
orange = 53-token chunk prefill, encoder co-scheduled in the SAME step (152/152 in 20 ticks) · blue = one decode token



Three ticks: ~870 ms busy / ~1130 ms idle per 2 s tick — the idle stripe is the conveyor's transfer window

