


Who is the engine actually serving? — per-session service timeline, vanilla vLLM-realtime, N=8, steady state ($t \approx 200$ s)

Requested 2-tick (4 s) zoom — audio for ALL 8 sessions arrives at the same instant (dotted = batched gateway push);
the engine is inside ONE session's service burst at a time (each block $\approx$ one 2 s chunk: encoder + prefill 53 tok + decode quota)



One full rotation (~15 s): emergent round-robin — each session gets a ≈ 1.9 s solo burst draining ≈ 7 backlogged chunks,
then waits ≈ 13 s for its next turn. No designed time-phase anywhere: this schedule is emergent, uncontrolled, and $7\times$ coarser than the tick

